


- NN_{init} weights & biases = 0.
 - Each neuron in the 1st h.l. will perform the same computation. Thus, even after multiple GD iterations, each neuron will be computing the same thing as other neurons.
- Log. reg. doesn't have a hidden layer. If you init. the weights to 0, the first ex. x fed will output 0 but the derivatives of the log. reg. depend on input x (because no hidden layer), and $x \neq 0$. Thus, at the 2nd iter, the weights follow x 's distribution and are different from each other. x isn't a const. vector. So the log. reg. won't fail to "break symmetry" if we init to all 0s.
- ~~NN~~ NN w/ tanh; init. w's ~~are~~ using $\text{random.randn}()$.
 - Inputs of tanh become very large, so gradients ≈ 0 , ~~slow~~ \therefore slow optimisation



$W^{[1]}$ is $(n^{[0]}, n^{[1]})$
 $b^{[1]}$ is $(n^{[0]}, 1)$

$W^{[1]}$ has shape $(n^{[0]}, n^{[1]})$ or $(4, 2)$	$Z^{[1]}$ $(4, m)$	\nearrow h.n.
$b^{[1]}$ $(n^{[0]}, 1)$ $(4, 1)$	$A^{[1]}$ $(4, m)$	\nearrow t.c.
$W^{[2]}$ $(n^{[1]}, n^{[2]})$ $(1, 4)$		
$b^{[2]}$ $(n^{[1]}, 1)$ $(1, 1)$		