Programming Assignment Overview: Planar Data Classification with one Hidden Layer

Learning objectives:
- Implement a 2-class classification NN w/ 1 HL
- Use units w/ a non-linear AF (e.g. tanh)
- Compute cross entropy loss
- Implement for. & backprop.

Contents:
1. Packages
2. Dataset
3. Simple log. reg.
4. NN model
   - 4.1. Defining the NN structure
   - 4.2. Initialise the model's parameters
   - 4.3. Loop

FP: Example $x^{(i)}$:

$$z^{[1](i)} = W^{[1]} x^{(i)} + b^{[1]}$$
$$a^{[1](i)} = \tanh(z^{[1](i)})$$
$$z^{[2](i)} = W^{[2]} a^{[1](i)} + b^{[2]}$$
$$\hat{y}^{(i)} = a^{[2](i)} = \sigma(z^{[2](i)})$$

→ np.dot(,)

$$y^{(i)}_{prediction} = \begin{cases} 1 & \text{if } a^{[2](i)} > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

Given predictions on all examples, cost:

$$J = -\frac{1}{m} \sum_{i=0}^{m} \left(y^{(i)} \log(a^{[2](i)}) + (1 - y^{(i)}) \log(1 - a^{[2](i)}) \right)$$

⋮

- 4.4. Integrate 4.1, 4.2, 4.3 in nn_model()
- 4.5. Predictions
- 4.6. Tuning hidden layer size
  - Larger models (more HUs) can fit the train. set better, until they overfit
  - n_h=5 seems to be the best

5. Performance on other datasets